

# Making the Web Conversational

How protocols like HTTP have evolved with modern Web applications

**Vipul Mathur**

November 8, 2017

*This talk is non-normative ;)*

Not speaking on behalf of my employer

# Overview

- Why
  - Life of the Web so far
  - Forms of communication
- What
  - Plain HTTP recap
  - Long polling and streaming
  - The WebSocket protocol
- How
  - WebSocket demo
  - Popular bi-directional Web frameworks
- Q&A

# Life of The Web: Protocols, Standards, Apps

## 1990: Cute Baby

- HTTP 0.9
- HTTP/1.0

## 2000: Responsive Teenager

- CSS2
- HTTPS popular

## 2010: Social Adult

- WebSocket
- SPDY
- WhatsApp

## 2020: Intelligent Elder?

## 1995: Interactive Child

- HTTP/1.1
- JavaScript (Browser)

## 2005: Chatty Adolescent

- JavaScript (Server)
- AJAX
- Gmail
- Google Docs

## 2015: Collaborative Worker


- HTTP/2
- IoT

# Forms of Communication

1. **Unicast** (1 to 1, simplex)
    - letters (snail-mail)
  2. **Multicast** (1 to many, simplex)
    - email
  3. **Broadcast** (1 to any, simplex)
    - TV/ radio
  4. **Request-response** (1 to 1, half duplex)
    - Q&A like in an interview
  5. **Conversation** (1 to 1, full duplex)
    - chat between friends
  6. **Omni-directional** (many to many, full duplex)
    - group chat or honking in Bangalore traffic :)
- Web protocols are mostly 4, increasingly 5
  - Web applications are of all types, built on top of Web protocols


# Plain HTTP Recap

- Request-response based
- Plain-text headers and body
  - human and machine readable



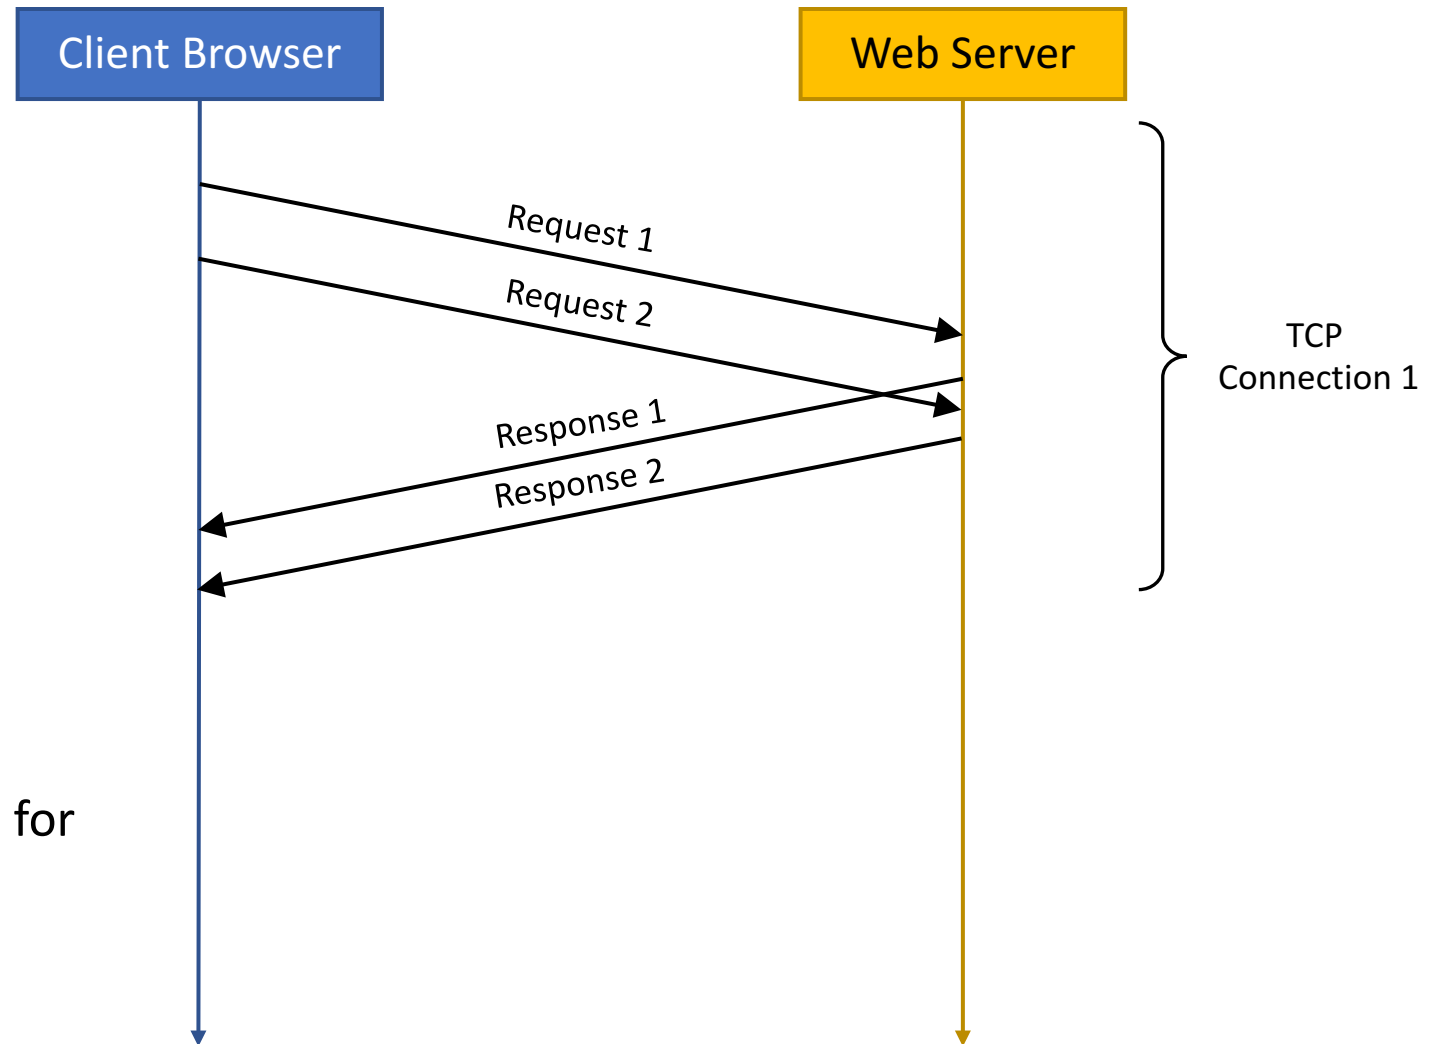
# Plain HTTP Recap

- Request-response based
- Plain-text headers and body
  - human and machine readable
- Persistent connections
  - multiple request-response pairs
  - on single TCP connection



# Plain HTTP Recap


- Request-response based
- Plain-text headers and body
  - human and machine readable
- Persistent connections
  - multiple request-response pairs
  - on single TCP connection
- Pipelining
  - multiple requests without waiting for response
  - on single TCP connection





# Plain HTTP Recap

- Request-response based
- Plain-text headers and body
  - human and machine readable
- Persistent connections
  - multiple request-response pairs
  - on single TCP connection
- Pipelining
  - multiple requests without waiting for response
  - on single TCP connection
- Chunked transfer encoding
  - allow response to be broken into chunks
  - allow headers after body



# Bi-directional Communication Over HTTP

## **Client to Server (TCP Conn #1)**

- Regular HTTP requests

## **Server to Client (TCP Conn #2)**

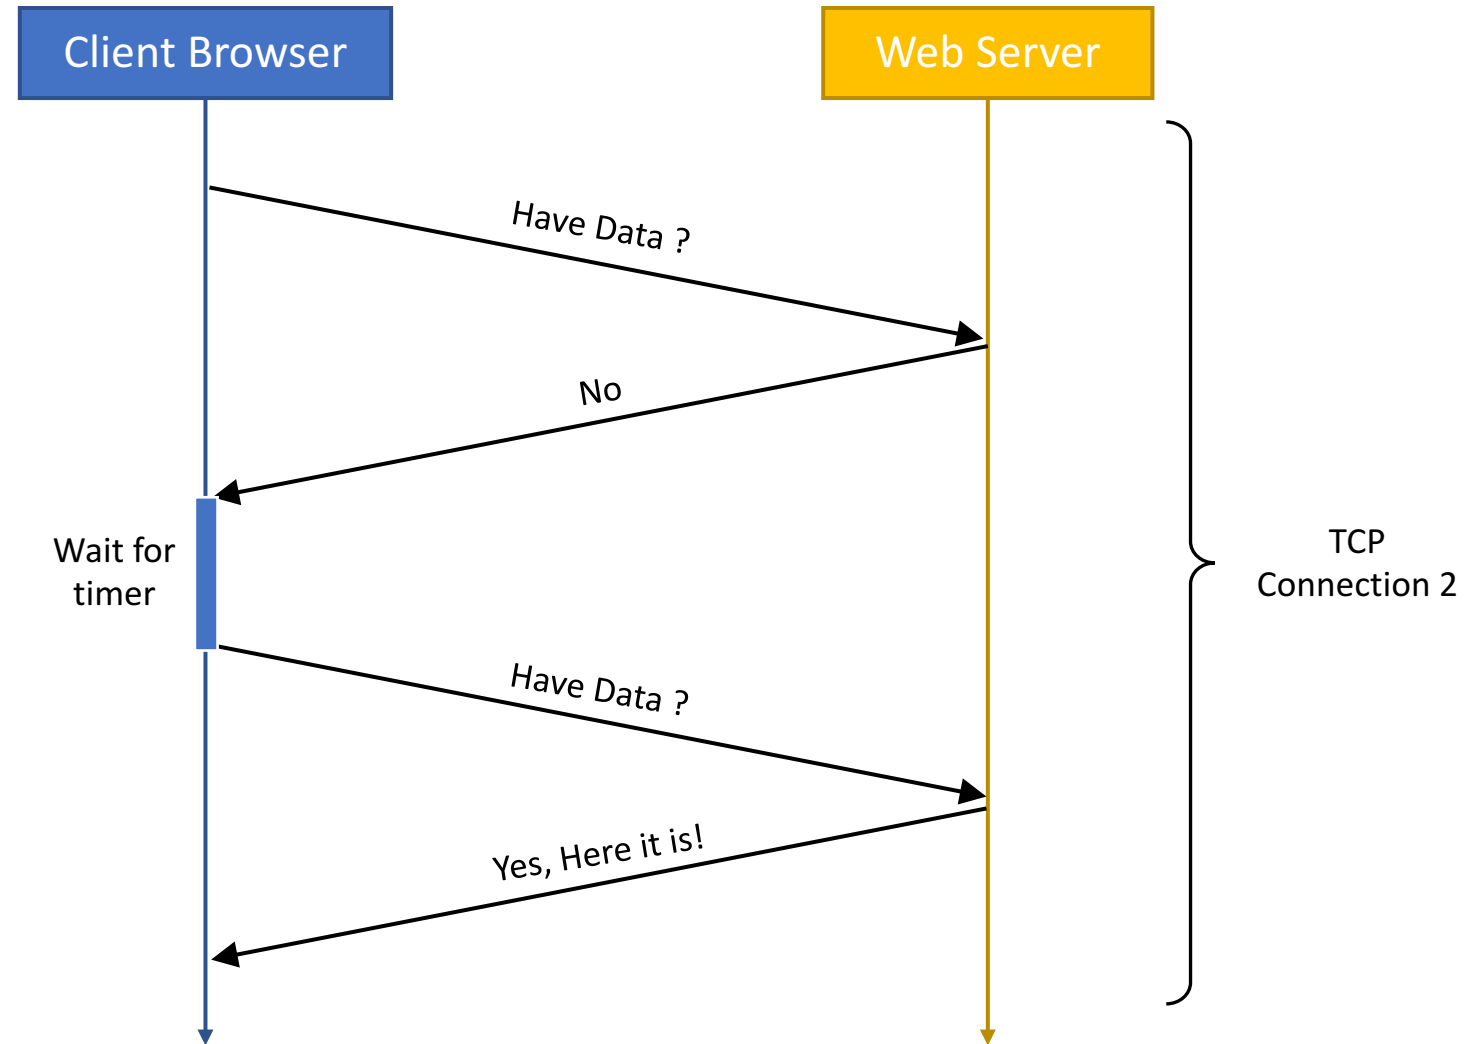
- Periodic Polling
- Long polling
- Streaming

## **Several Practical Combinations Implemented**

- BOSH (Bidirectional-streams Over Synchronous HTTP)
- Comet (e.g. Pushlets)
- Bayou
- Server-Sent Events

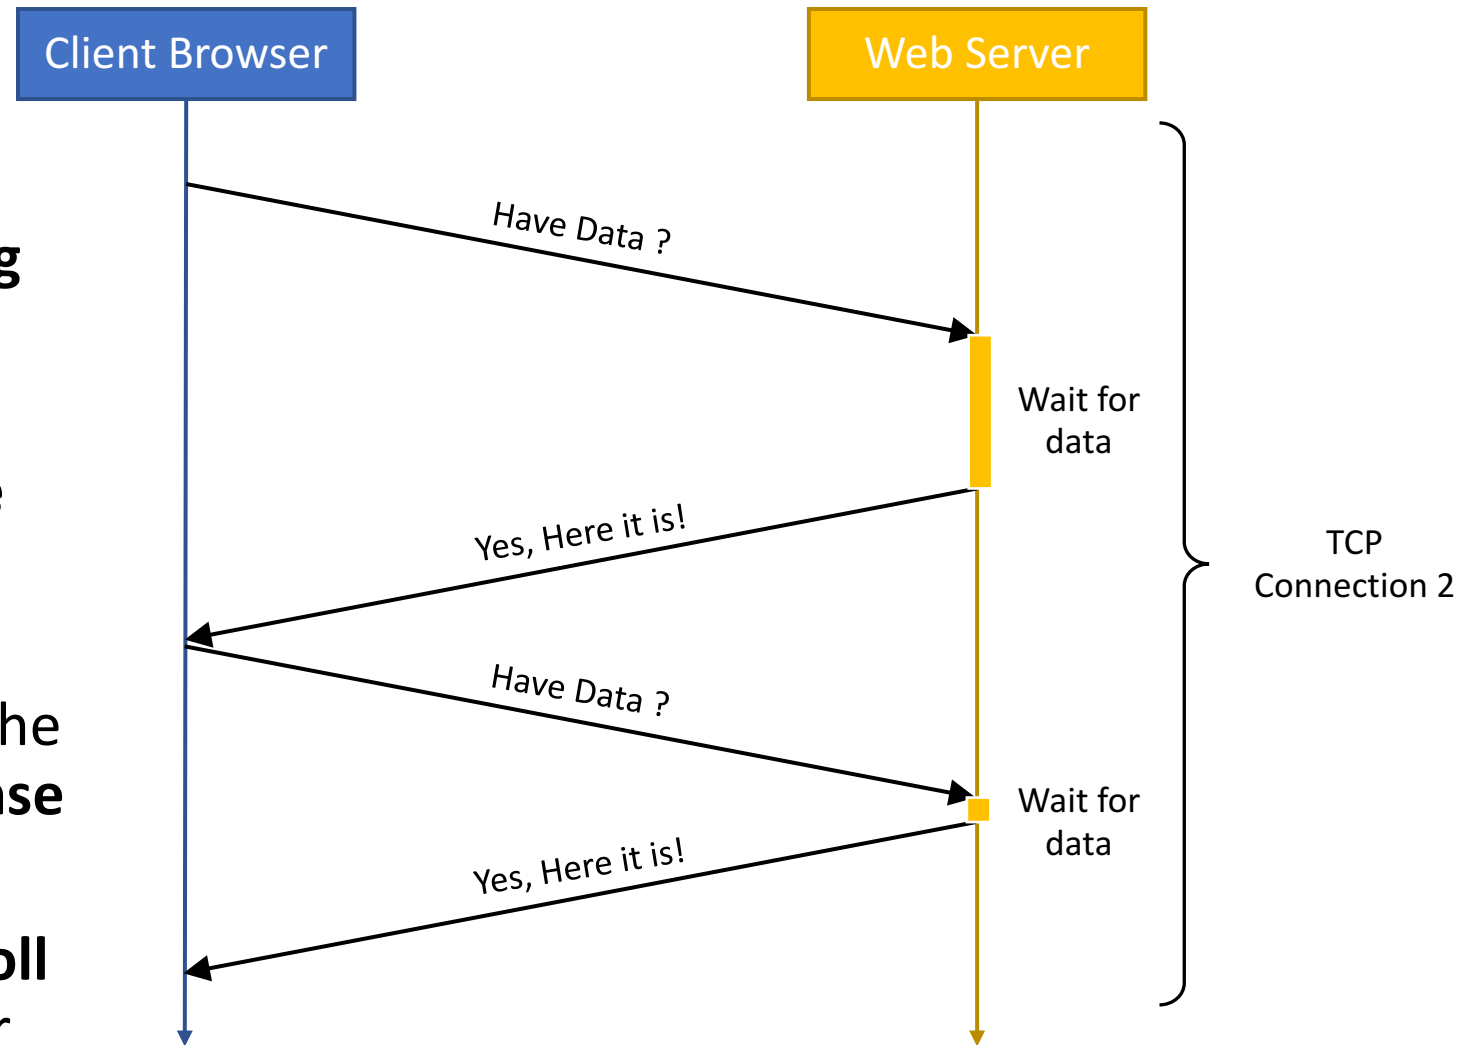
# Periodic Polling

1. The client makes a **polling request** to check if server has data.
2. The **server sends a response** to the client, even if there is no data.
3. The client waits for some time and **repeats polling** from step 1.



# HTTP Long Polling

1. The client makes an **initial long poll request** and waits for a response.
2. The **server defers its response** until an update is available, or timeout has occurred.
3. When an update is available, the **server sends complete response** to the client.
4. The client sends a **new long poll request**, either immediately or after a pause.



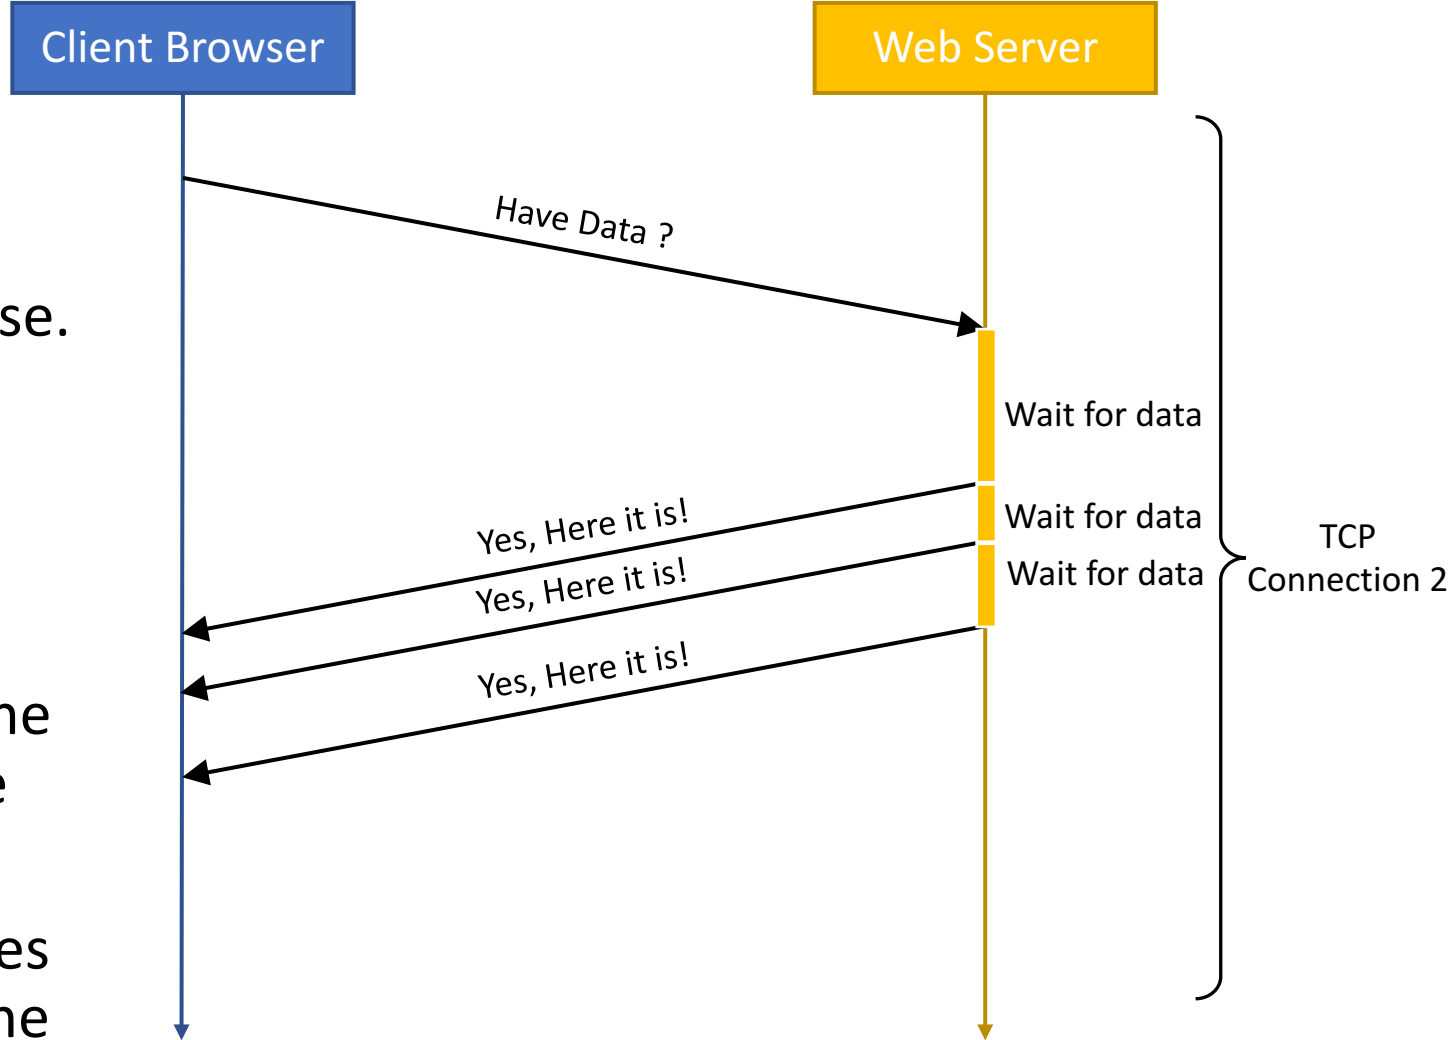
# Long Polling: Issues

## Issues

1. Header overhead
2. Unnecessary high maximal latency ( $> 1.5$  RTT)
3. Frequent TCP connections (persistence helps)
4. Resource overhead at client and server
5. Unneeded buffering during higher loads
6. Timeouts
7. Caching

# HTTP Streaming

1. The client makes an **initial request** and waits for a response.
2. The **server defers its response** until an update is available, or until a particular status or timeout has occurred.
3. When an update is available, the **server sends a response** to the client.
4. The data sent by the server does not terminate the request or the connection. The **server returns to step 3**.



# Streaming: Issues

1. Network intermediaries (proxies, gateways) may buffer response
2. Unnecessary maximal latency ( $> 1.5$  RTT) due to re-establishing streaming to avoid client memory limits
3. Client buffering (library to app)
4. Framing requirements not met by chunking (due to re-chunking)

# Pushlets

- Old ~2002
- Implementation of Comet
- Publish/ subscribe mechanism
- Java servlets on server
- Push JavaScript snippets from server to client using HTTP streaming
- [Whitepaper](#)



# Server-Sent Events (SSE)

- Proposed ~2009
- W3C Recommendation ~2015
- JavaScript [EventSource API](#) part of HTML5
- Built over HTTP streaming
- Good support by most modern browsers
  - Not supported by IE
  - 'Under Consideration' for Edge 16
- Uses Content-Type: text/event-stream

# Upgrade to WebSocket!

RFC 6455

“Historically, creating web applications that need bidirectional communication between a client and a server (e.g., instant messaging and gaming applications) has required an **abuse of HTTP** to poll the server for updates while sending upstream notifications as distinct HTTP calls”

*Introduction of RFC6455 (WebSocket Protocol) referring to RFC6202 (Long Polling and Streaming Issues and Best Practices)*

(**emphasis mine**)

# WebSocket vs. Long Polling/ Streaming

## **Problems Addressed**

- Server resources: multiple TCP connections per client (up/ down)
- Protocol overhead: long header per HTTP message
- Client complexity: track pair of connections and state to one server

## **Solution Approach**

- Use a single TCP connection for traffic in both directions.

# WebSocket Protocol and API

- WebSocket protocol
  - Defined in [RFC6455](#)
  - Status: Proposed Standard ~2011
- JavaScript [WebSocket API](#)
  - Part of HTML5
  - W3C Candidate Recommendation ~2012
- Pointer: Watch the excellent talk [Inside WebSockets](#) by Leah Hanson

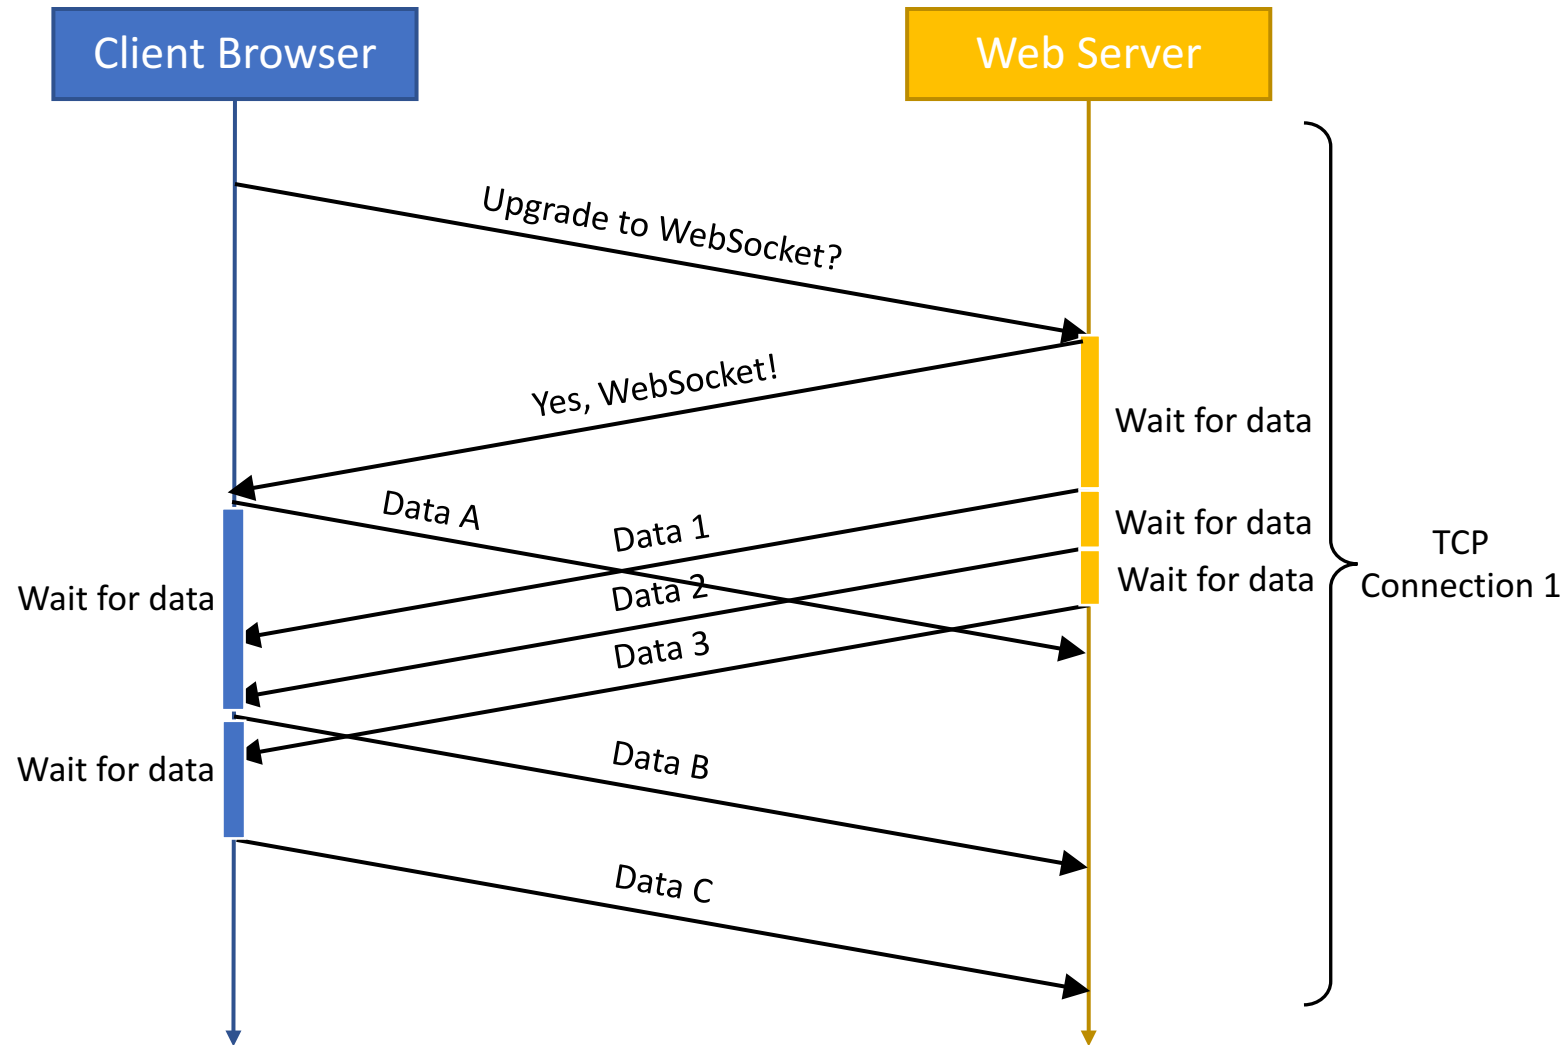
# WebSocket

## Advantage!

Only one TCP connection for both directions

## Basic Steps

1. Opening handshake
2. Data exchange (two-way)
3. Closing handshake



# WebSocket: Highlights

- Bi-directional communication over single TCP connection
- Either client or server can send a message anytime
- Low resource overhead at client and server
- No maximal RTT (except for opening handshake)
- Higher efficiency due to binary framing
- Designed to work well with existing Web infrastructure
  - Start with plain HTTP and 'upgrade' to WebSocket
  - Uses ports 80 and 443 for WS and WSS
  - Can tunnel through HTTP proxies via HTTP CONNECT

# WebSocket Demo

- Simple echo client
  - <https://www.websocket.org/echo.html>
  - <http://janodvarko.cz/test/websockets/>
- Inspect WS with Firefox/ Chrome developer tools
  - Ctrl/Cmd + Shift + I
  - Network Tab → WS
- Firefox [WebSocket Monitor](#) extension
  - Extend developer tools with better WS support
  - Visualize WS sessions
  - <https://github.com/firebug/websocket-monitor>



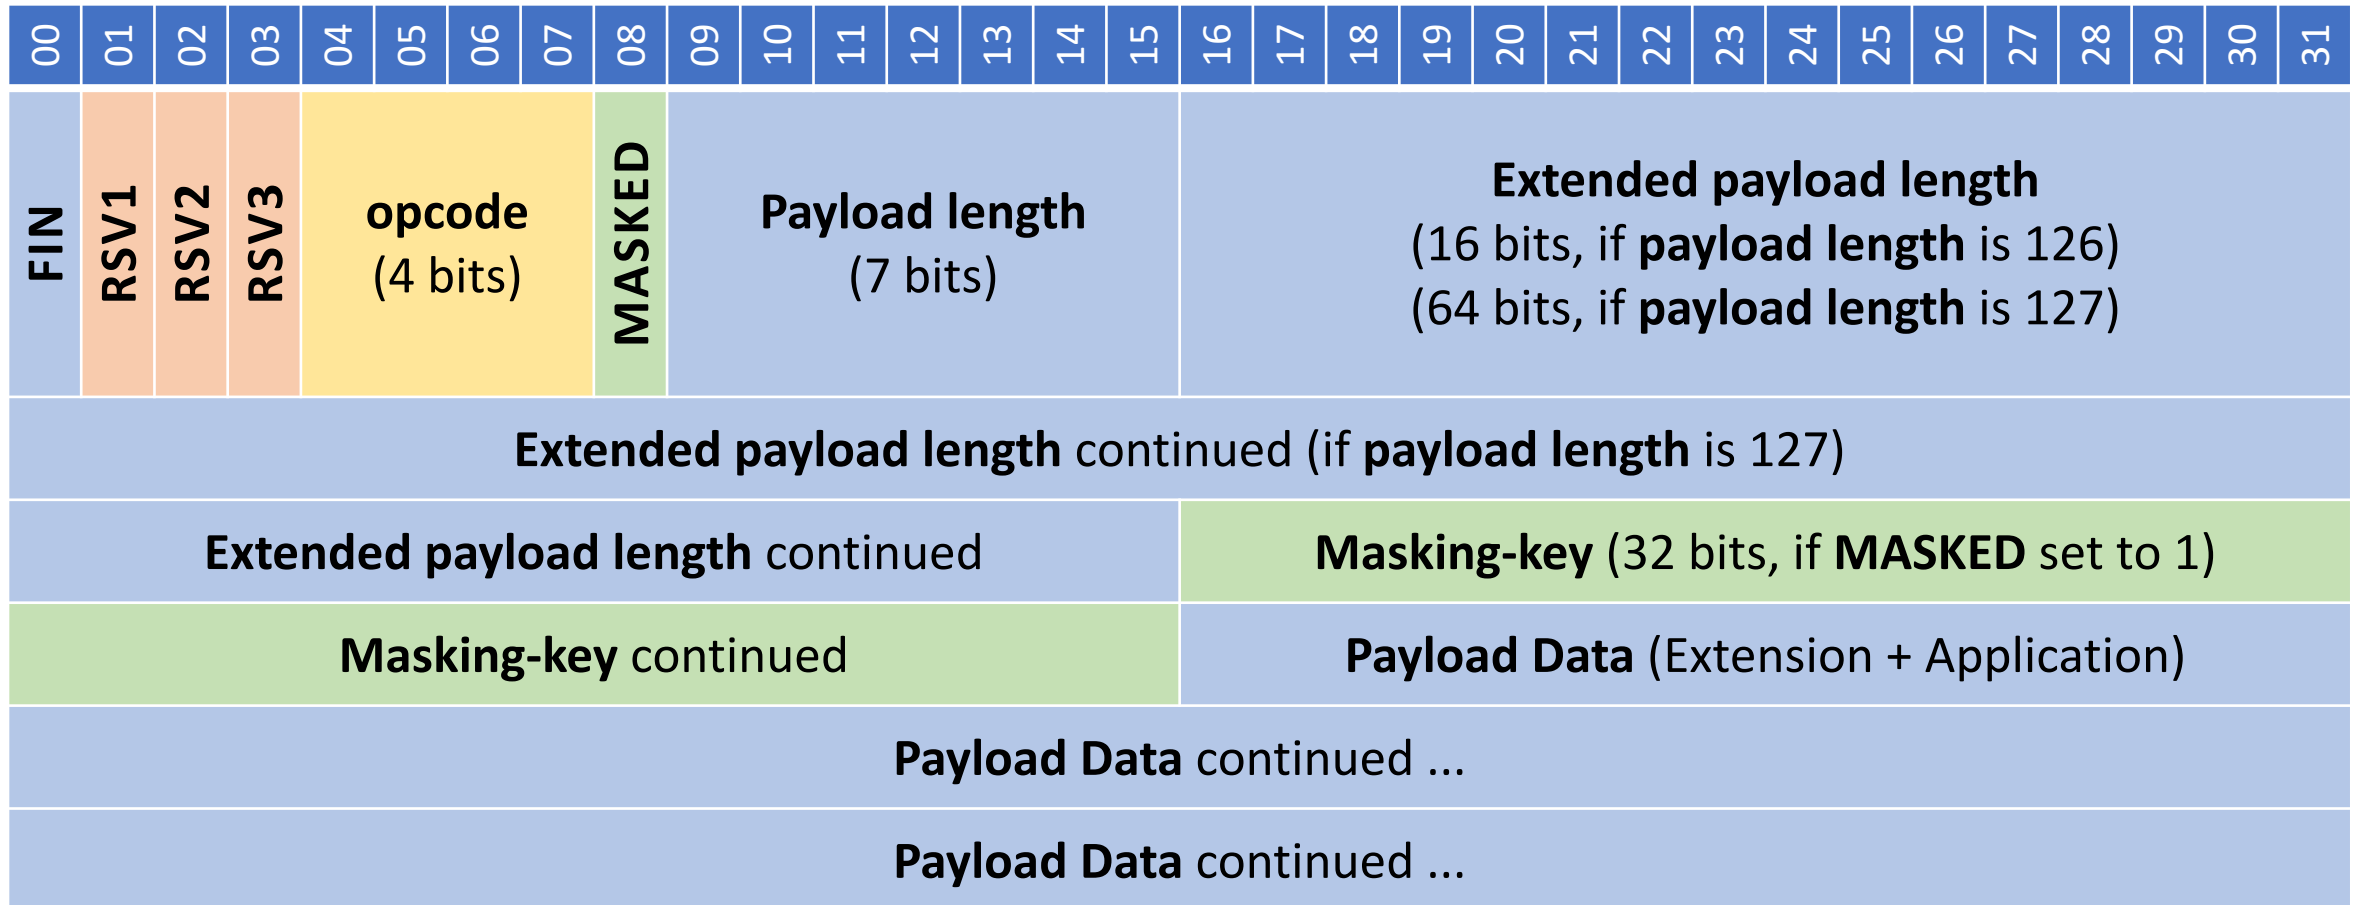
# WS: Other Notable Points

- Text (UTF-8) and binary data
  - Messages and Framing
  - Ping, Pong for keep-alive
  - Frame masking from client to server
- 
- Good implementations in browsers, and support in servers
  - WS plain and WSS using TLS
  - Sub-protocols (e.g. chat)
  - Extensions (e.g. compression)

# WebSocket Data Framing

00	01	02	03	04	05	06	07	08	09	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

# WebSocket Data Framing



# WebSocket Extensions

## **Multiplexing (HyBi WG Draft)**

- Provides logical channels in one TCP connection
- Each channel equivalent to a WebSocket connection
- Avoid head-of-line blocking by using different channels
- Current state: Expired (~2014)
- HTTP/2 maybe?

## **Compression (RFC 7692)**

- Compress payload data of a message
- Negotiate parameters during opening handshake
- DEFLATE algorithm default
- Current state: Proposed standard (~2015)
- Browser support: minimal

# Interesting Bi-Directional Frameworks

## SockJS: WebSocket emulation

- WebSocket-like API even without WebSocket transport
- Focus on cross-browser compatibility
- Try native WebSocket first
- Automatically fall back transports (WS > streaming > polling)
- Multiple server implementations (JavaScript, Python, Java, Scala, Ruby, Go, Erlang, ...)

## Socket.IO: Real-time framework

- HTTP long-poll first
- Later upgrade to WebSocket if possible
- Handles disconnects
- Built-in keep-alive
- Supports namespaces
- Goes well beyond just WebSocket semantics
- JavaScript only

# Q&A

## **My Coordinates**

Twitter: [@VipulMathur](https://twitter.com/VipulMathur)

LinkedIn: <https://linkedin.com/in/VipulMathur>